

CS3114 (Fall 2016)

PROGRAMMING ASSIGNMENT #2

Due Wednesday, October 12 @ 11:00 PM for 100 points

Due Tuesday, October 11 @ 11:00 PM for 10 point bonus

Note: This project also has three intermediate milestones. See the Piazza forum for details.

Last updated: 10/3 at 10:00 am

Assignment:

This project builds on Project 1. In Project 1, the hash table was great for keeping you from storing into the memory pool multiple copies of a given artist name or a given song name. But there was no support to answer questions like “What are all of the songs by artist X?” or “What artists recorded the song (with name) Y?”. In this project you will add a data structure (called a 2-3⁺ tree) that supports **range queries**, so that these questions can be answered. The hash tables and the memory manager will be almost unchanged from Project 1.

Invocation and I/O Files:

The program would be invoked from the command-line as:

```
java SearchTree {initial-hash-size} {block-size} {command-file}
```

The name of the program is **SearchTree**. Parameter **{initial-hash-size}** is the initial size of the hash table (in terms of slots). Parameter **{block-size}** is the initial size of the memory pool (in bytes). The meaning of the parameters is identical to Project 1.

Your program will read from text file **{command-file}** a series of commands, with one command per line. The program should terminate after reading the end of the file. The formats for all commands are identical to Project 1. However, these new commands are added.

```
list {artist|song} {name}
```

If the first parameter is **artist** then all songs by the artist with name **{name}** are listed. If the first parameter is **song** then all artists who have recorded that song are listed. They should be listed in the order that they appear in the 2-3⁺ tree.

```
delete {artist-name}<SEP>{song-name}
```

Delete the specific record for this particular song by this particular artist. This means removing two records from the 2-3⁺ tree, undoing the insert. If this is the last instance of that artist or of that song, then it needs to be removed from its hash table and the memory pool. In contrast, the **remove** command removes all instances of a given artist or song from the 2-3⁺ tree as well as the hash table and memory pool.

```
print tree
```

You will print out a pre-order traversal of the 2-3⁺ tree as follows. First, before printing the tree, you should print out on its own line: “**Printing 2-3 tree:**”. Each node is printed on its own line. Each node is indented by its depth times two spaces. So the root is not indented, its children are indented 2 spaces, its grandchildren 4 spaces, and so on. Internal nodes list the (one or two) placeholder records (both the key and the value for each one). That is, if an internal node stores a key (a handle) whose corresponding string is at position 10 in the memory pool, and a value handle with memory pool position 20, then you would print 10 20. Leaf nodes print, for each

record it stores, the two handle positions. So they will print either 2 or 4 numbers. The numbers should be separated by a space.

The 2-3⁺ Tree

You will implement a 2-3⁺ tree to support the search query. The data records to be stored are objects of the class `KVPair`, where the key is a handle, and the value is another handle. Your `Handle` class should be made to support the `Comparable` interface, and the actual thing that will be compared is the start position for the string in the memory manager's array. In other words, ultimately, the keys for the strings are their locations in the memory manager's array. However, to simplify implementation, we want to avoid storing records with duplicate keys. So while we use the "key" field when comparing records, if they have the same value in the key (because they refer to the same artist, or to the same song), we will then use the value of the handle in the value position of the `KVPair` to break the tie. For example, if one record has handle 10 in the key field and handle 30 in the value field, while another record has handle 10 in the key field and handle 40 in the value field, then the first one would appear in list of leaf nodes before the second.

Whenever an insert command is called with a song/artist pair, the handles for these two strings will then be used to enter **two** new entries into the 2-3⁺ tree. If the artist's name handle is named `artistHandle` and the song's name handle is named `songHandle`, then you will create and insert a `KVPair` object with the `artistHandle` as the key and the `songHandle` as the value field. You will then create and insert a second `KVPair` object with the `songHandle` as the key and the `artistHandle` as the value field. However, be sure not to insert a duplicate artist/song record into the database. If there already exists a `KVPair` object in the 2-3⁺ tree with those same handles, then you will not add the song pair again.

When a `list` command is processed, you will first get the handle for the associated string from the hash table. You will then search for that handle in the 2-3⁺ tree and print out all records that have been found with that key value.

When a `remove` command is processed, you will be removing all records associated with a given artist or song name. To (greatly!) simplify removing the records from the 2-3⁺ tree, you will remove them one at a time. So, you will search for the first record matching the corresponding handle for the name that you want to delete, then call the (single record) delete operation on it, and repeat this process for as long as there is a record matching that handle in the tree. Also, whenever you remove a particular record (say, an artist/song pair), you will immediately after remove the corresponding song/artist record. Finally, whenever you remove the last instance of either an artist or a song, you always should remove the associated record from the hash table and the string from the memory pool.

When implementing the 2-3⁺ tree, all major functions (insert, delete, search) must be implemented recursively. You may **not** store a parent pointer for the nodes. Your nodes must be implemented as a class hierarchy with a node base class along with leaf and internal node subclasses. Internal nodes store two values (of type `KVPair`) and three pointers to the node base class. Leaf nodes store two `KVPair` records.

Programming Standards:

You must conform to good programming/documentation standards. Web-CAT will provide feedback on its evaluation of your coding style, and be used for style grading. Beyond meeting Web-CAT's checkstyle requirements, here are some additional requirements regarding programming standards.

- You should include a comment explaining the purpose of every variable or named constant you use in your program.
- You should use meaningful identifier names that suggest the meaning or purpose of the constant, variable, function, etc. Use a consistent convention for how identifier names appear, such as “camel casing”.
- Always use named constants or enumerated types instead of literal constants in the code.
- Source files should be under 600 lines.
- There should be a single class in each source file. You can make an exception for small inner classes (less than 100 lines including comments) if the total file length is less than 600 lines.

We can’t help you with your code unless we can understand it. Therefore, you should no bring your code to the GTAs or the instructors for debugging help unless it is properly documented and exhibits good programming style. Be sure to begin your internal documentation right from the start.

You may only use code you have written, either specifically for this project or for earlier programs, or code provided by the instructor. Note that the OpenDSA code is not designed for the specific purpose of this assignment, and is therefore likely to require modification. It may, however, provide a useful starting point.

Deliverables:

You will implement your project using Eclipse, and you will submit your project using the Eclipse plugin to Web-CAT. Links to the Web-CAT client are posted at the class website. If you make multiple submissions, only your last submission will be evaluated. There is no limit to the number of submissions that you may make.

You are required to submit your own test cases with your program, and part of your grade will be determined by how well your test cases test your program, as defined by Web-CAT’s evaluation of code coverage. Of course, your program must pass your own test cases. Part of your grade will also be determined by test cases that are provided by the graders. Web-CAT will report to you which test files have passed correctly, and which have not. Note that you will **not** be given a copy of these test files, only a brief description of what each accomplished in order to guide your own testing process in case you did not pass one of our tests.

When structuring the source files of your project, use a flat directory structure; that is, your source files will all be contained in the project “src” directory. Any subdirectories in the project will be ignored.

You are permitted to work with a partner on this project. When you work with a partner, then **only one member of the pair** will make a submission. Be sure both names are included in the documentation. Whatever is the final submission from either of the pair members is what we will grade unless you arrange otherwise with the GTA.

Pledge:

Your project submission must include a statement, pledging your conformance to the Honor Code requirements for this course. Specifically, you must include the following pledge statement near the beginning of the file containing the function `main()` in your program. The text of the pledge will also be posted online.

```
// On my honor:
//
// - I have not used source code obtained from another student,
//   or any other unauthorized source, either modified or
//   unmodified.
//
// - All source code and documentation used in my program is
//   either my original work, or was derived by me from the
//   source code published in the textbook for this course.
//
// - I have not discussed coding details about this project with
//   anyone other than my partner (in the case of a joint
//   submission), instructor, ACM/UPE tutors or the TAs assigned
//   to this course. I understand that I may discuss the concepts
//   of this program with other students, and that another student
//   may help me debug my program so long as neither of us writes
//   anything during the discussion or modifies any computer file
//   during the discussion. I have violated neither the spirit nor
//   letter of this restriction.
```

Programs that do not contain this pledge will not be graded.